

Engineering Adaptive Software Systems: Policies

Marin Litoiu

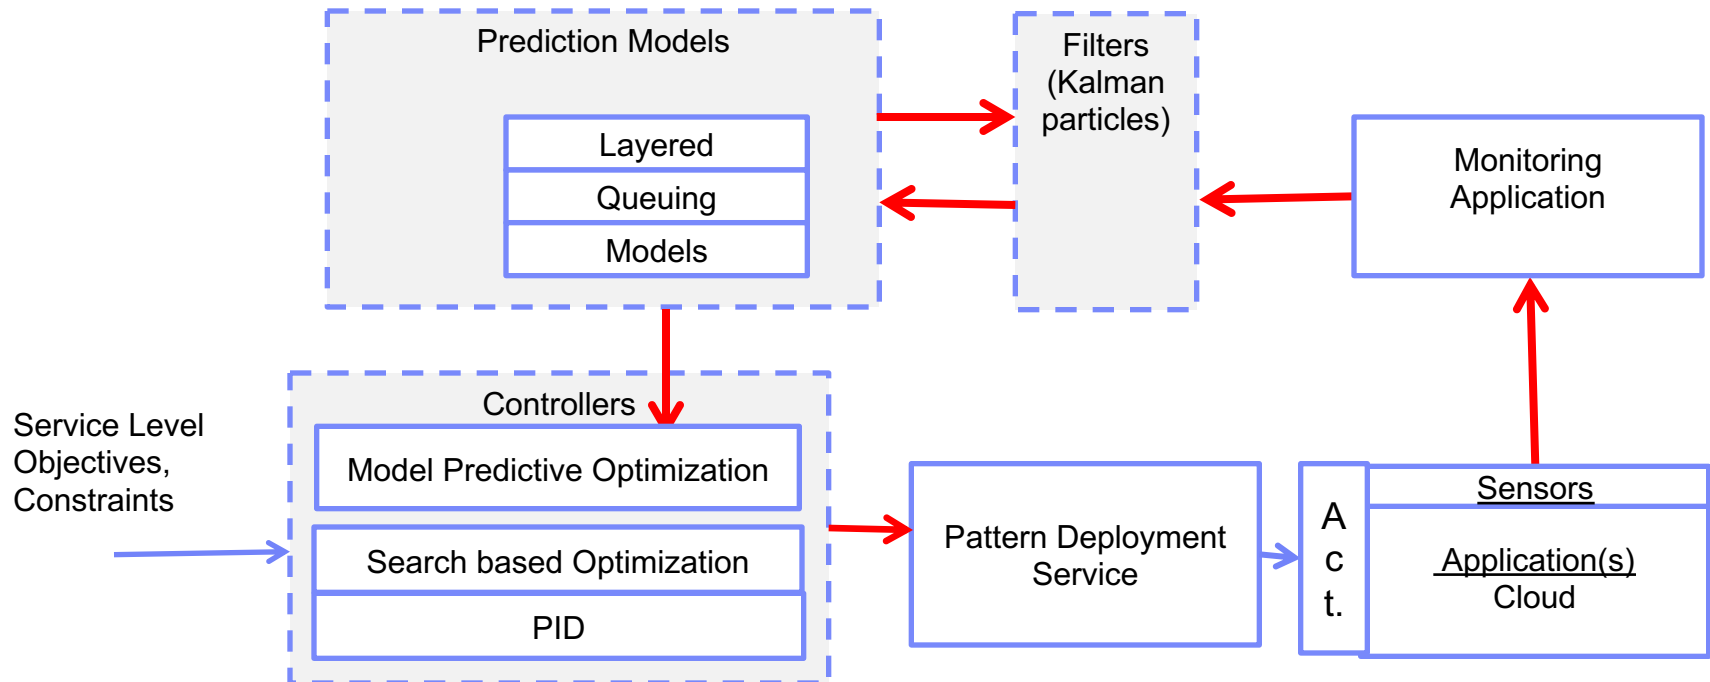
Adaptive Systems Research Lab

York University

Department of EECS and School of IT

<http://ceraslabs.com>

Look Ahead Adaptation



- **Predictive (red arrows): anticipates future load, performance, cost**
 - Uses prediction models, filters and predictive optimization. It is slow and effortful but efficient

Look Ahead Optimization with Thrashing in Cloud

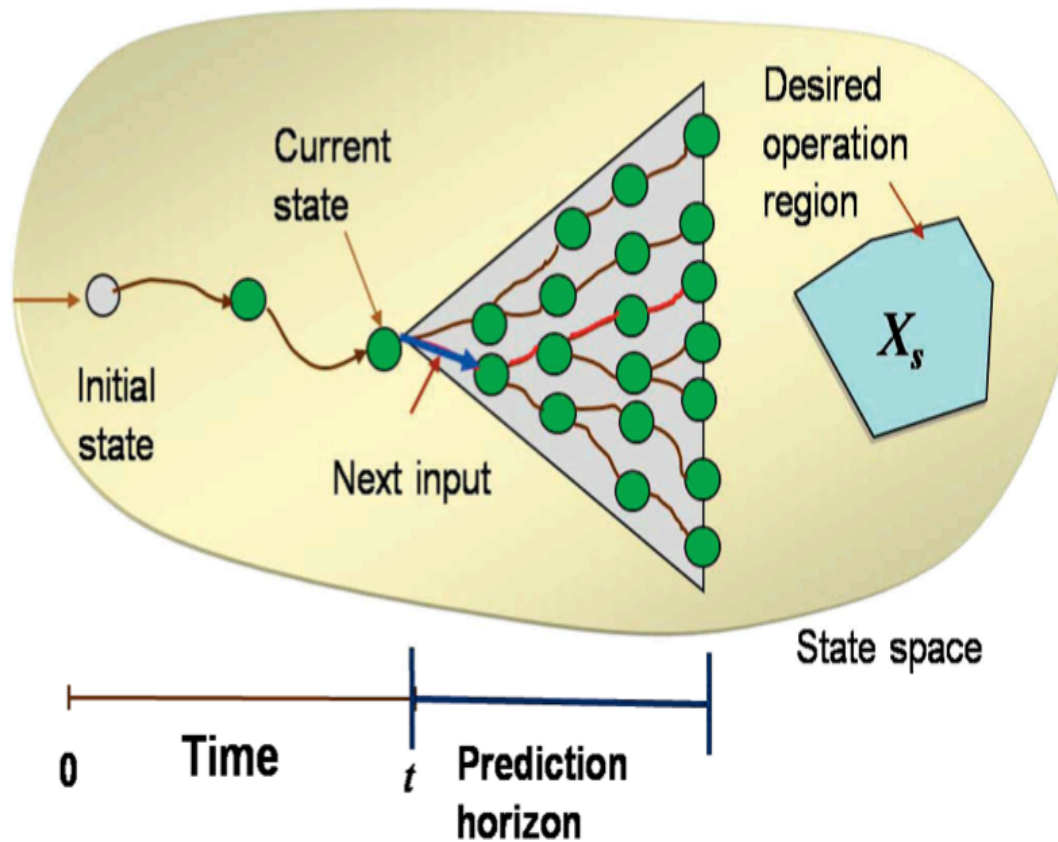
- **The problem:**

- Given a set of applications in a private cloud, with variable workloads
 - Minimize the cost, performance, thrashing
 - By allocating/migrating VMs to physical machines in cloud

- **How it works:**

- Monitor the current states
- Predict the workload over a look ahead horizon T
- Compute the allocation for each step $t=1, 2 \dots T$
 - For each step, use the model of the system to predict the cost
- Implement the allocation for step 1
- Repeat Monitor, Predict, Compute, Implement steps....

Look Ahead Adaptation



From [Bai and Abdelwahed]

Self-Managing Policies

- Autonomic elements function at different levels of abstraction
- At the lowest levels, the capabilities and the interaction range of an autonomic element are limited and hard-coded
- At higher levels, elements pursue more flexible goals specified with policies, and the relationships among elements are flexible and may evolve over time
- Action, goal and utility-function policies

Kephart & Walsh; An AI Perspective on AC Policies, 5th IEEE International Workshop on Policies for Distributed

Policy Examples

- **A policy is a set of considerations designed to guide decisions of courses of action.**
- **“Neither a borrower, nor a lender be; for a loan oft loses both itself and friend, and borrowing dulls the edge of husbandry.” In *Hamlet*, Shakespeare’s policy regarding borrowing.**
- **Star Wars**
 - When C3PO, upon receiving caution from Hans Solo, tells R2D2 to “let the wookie win.” Apparently Chewbacca (the wookie in question) had a habit of detaching an opponent's arm upon losing.
 - It is important to note that R2D2 had another implicit policy that said when he's competing, he should try to win, and this policy directly conflicted with Solo's sage advice.
 - In the end, R2D2 let the Wookie have the game, valuing his arm over the victory.

Action Policies


- **Dictate the actions that should be taken when the system is in a given state**
- **IF (condition) THEN (action)**
 - where the condition specifies either a specific state or a set of possible states that all satisfy the given condition
- **Note that the state that will be reached by taking the given action is not specified explicitly**
- **Policy author knows which state will be reached upon taking the recommended action and deems this state more desirable than states that would be reached via alternative actions**

Goal Policies

- Rather than specifying exactly what to do in the current state, goal policies specify either a single desired state, or one or more criteria that characterize an entire set of desired states
- Rather than relying on a human to explicitly encode rational behavior, as in action policies, the system generates rational behavior itself from the goal policy
- This type of policy permits greater flexibility and frees human policy makers from the “need to know” low-level details of system function, at the cost of requiring reasonably sophisticated planning or modeling algorithms

Utility-Function Policies

- An objective function that expresses the value of each possible state
- Generalized goal policies
- Instead of performing a binary classification into desirable versus undesirable states, they ascribe a real-valued scalar desirability to each state
- Because the most desired state is not specified in advance, it is computed on a recurrent basis by selecting the state that has the highest utility from the present collection of feasible states
- Provide more fine-grained and flexible specification of behavior than goal and action policies
- Allow for unambiguous, rational decision making by specifying the appropriate tradeoff
- Preferences are difficult to elicit and specify



Autonomic Computing Policies

- What is the difference between **action, goal and utility-function policies**?
 - Advantages, disadvantages, benefits, limitations?



Policy Types for Autonomic Computing

■ Action policies

- **If-then action rules** specify exactly what to do under the current condition.
- Rational behaviour is compiled in by the designer

■ Goal policies

- Requires self-model, planning, conceptual knowledge representation


■ Utility function policies

- It chooses the actions to maximize its utility function
- Finer distinction between desirability of different states than goals
- Numerical characterization of state
- Needs methods to carry out actions to optimize utility

Real autonomic systems embody a combination of policy types.

Action Policies

- A state S is a vector of attributes
- S can directly be measured by a sensor, or
- S can be inferred or synthesized from lower-level measurements
- Policy will directly or indirectly cause an action a
- Deterministic or probabilistic transition into a new state from S to a new state T




Action Policy Example

- **RT: Response Time**

if ($RT_{\text{Gold}} > 100 \text{ ms}$) increase CPU_{Gold} by 5%

if ($RT_{\text{Silver}} > 200 \text{ ms}$) increase CPU_{Silver} by 5%



Action Policy Example

- **For each machine, if idle session is greater than 20 minutes then terminate the session**

- **BitTorrent user processes initiated from IP address 141.223.2.15 should have lowest priority**

if (srcIPaddress == 141.223.2.15 && process-type == "bittorrent")

then priority is low

- **Event**

- Total number of user logins is greater than 5 **and**
- CPU load is greater than 90 **and**
- Total number of processes running is greater than 35

- **Condition**

- None

- **Action**

- Block any new user logins

Goal Policies

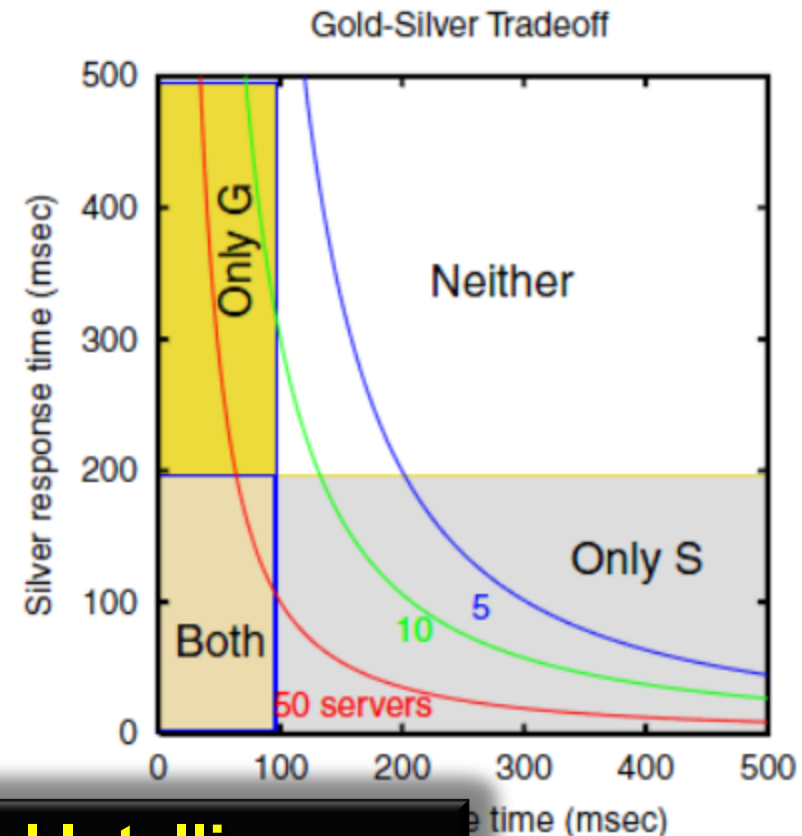
- Instead of specifying what to do in a current state, specify a single desired state or a set of desired states are
- Any member of this target set is equally acceptable
- Cannot express fine distinctions in preference
- How to compute a set of actions that gets the system from current state S to a desired state T ?
- The system generates rational behaviour

Goal Policy Example

- RT: Response Time

Gold: $RT_{\text{Gold}} \leq 100 \text{ ms}$

Silver: $RT_{\text{Silver}} \leq 200 \text{ ms}$



**Kephart and Walsh: An Artificial Intelligence
Perspective**

on Autonomic Computing Policies, POLICY

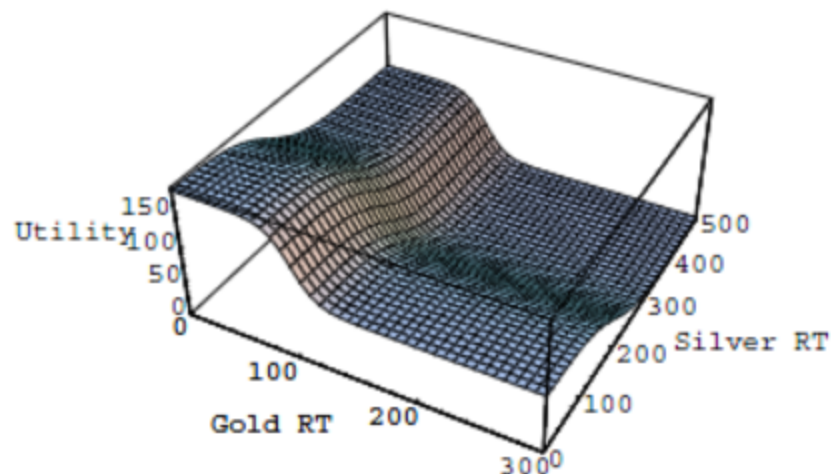
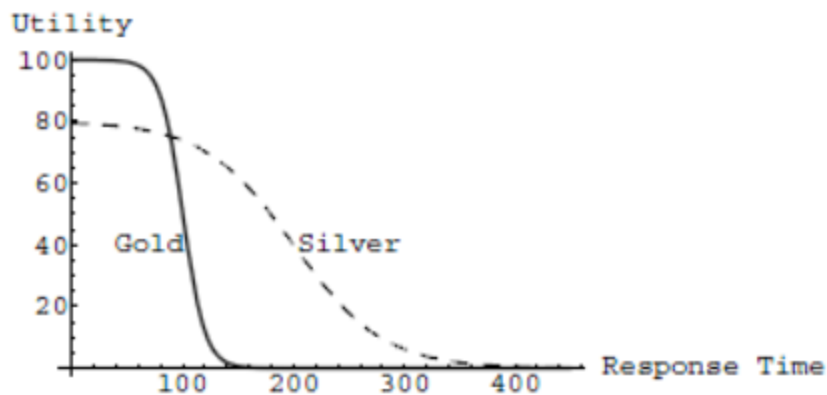
Utility-Function Policies

- An objective function to express the value of each possible state
- Generalizes goal policies
- Instead of desirable/undesirable we have a real-valued scalar desirability for each state
- More fine-grained and flexible specifications
- Goal functions often exhibit conflict
- Unambiguous, rational decision making

Utility-Function Policy Example

- **U: Utility function**

$$U(RT_{\text{Gold}}, RT_{\text{Silver}}) = U_{\text{Gold}}(RT_{\text{Gold}}) + U_{\text{Silver}}(RT_{\text{Silver}})$$



Policy Types for Autonomic Computing

- **Real autonomic systems embody a combination of policy types**
- **In ACRA**
 - Lower levels typically use action policies
 - For higher levels, goal or utility-function policies are more appropriate
- **Unified framework is needed to support multiple policy types within a single autonomic component**

Policy Languages

- **PDL**
- **Ponder**
- **Rei**
- **KAoS**
- **ACPL, SPL (PMAC)**
- **XACML**
- **CIM-SPL**

Ponder

- **Developed @ Imperial College London**
 - A declarative, object-oriented language for specifying different types of policies
 - grouping policies into roles and relationships
 - defining configurations of roles and relationships as management structures.
 - Can be used to specify security policies with role-based access control, as well as general purpose management policies

More info

**Kephart and Walsh: An Artificial Intelligence
Perspective
on Autonomic Computing Policies, POLICY
2004.**

TOWARDS A SE METHODOLOGY

Engineering Adaptive Loops: Methodology*

*Parisa Zoghi, Mark Shtern, Marin Litoiu, and Hamoun Ghanbari. 2016. Designing Adaptive Applications Deployed on Cloud Environments. *ACM Trans. Auton. Adapt. Syst.* 10, 4, Article 25

- **Elicitation phase (Goals and operations)**
- **Ranking phase**
- **Adaptation algorithm phase**

1. Adaptation Goals



2. Control Points

- **Definition**

- Operations that called at run-time cause controlled changes in a system
- Can belong to services, applications, cloud

- **Examples**

- Add/remove instances
- Launch instances in cheaper availability zone

3. Ranking Phase


Control point alternatives	Numerical values
Alternative 1	1
Alternative 2	2
...	...
Alternative n	n

4. Implement MAPE Loops

- **Monitor Goals**
- **Analyze and Plan**
 - Compute a Strategy over the space of Control Points using a Model
 - MPC
 - LQR
 - Search Based Algorithms
- **Execute the Strategy**

Conclusions

Self 1 and Self 2*



- Intuitive
- Reactive
- Fast
- Low resource consumption
- Unconscious
- Useful but not always right
- Ex: $2 \times 9 = ?$

- Analytical
- Deliberative
- Slow
- High resource consumption
- Conscious
- More often right
- Ex: $69 \times 37 = ?$

Reactive Adaptation



Model Based
Adaptation